

Machine learning for vision and multimedia

(01URPOV)

Lab 04 – Visualizing deep neural networks
Lia Morra

2024 – 2025

Visual analysis in deep learning

Visual Analytics in Deep Learning | Interrogative Survey Overview

§4 WHY

Why would one want to use visualization in deep learning?

Interpretability & Explainability
Debugging & Improving Models
Comparing & Selecting Models
Teaching Deep Learning Concepts

§6 WHAT

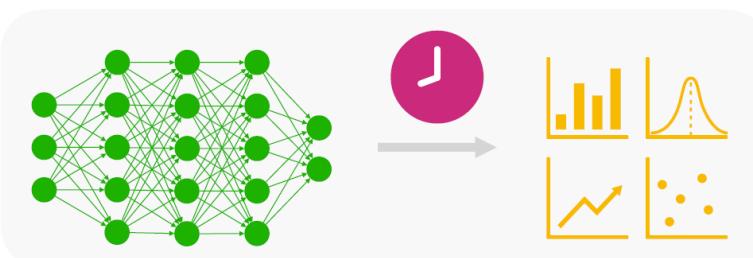
What data, features, and relationships in deep learning can be visualized?

Computational Graph & Network Architecture
Learned Model Parameters
Individual Computational Units
Neurons In High-dimensional Space
Aggregated Information

§8 WHEN

When in the deep learning process is visualization used?

During Training
After Training



§5 WHO

Who would use and benefit from visualizing deep learning?

Model Developers & Builders
Model Users
Non-experts

§7 HOW

How can we visualize deep learning data, features, and relationships?

Node-link Diagrams for Network Architecture
Dimensionality Reduction & Scatter Plots
Line Charts for Temporal Metrics
Instance-based Analysis & Exploration
Interactive Experimentation
Algorithms for Attribution & Feature Visualization

§9 WHERE

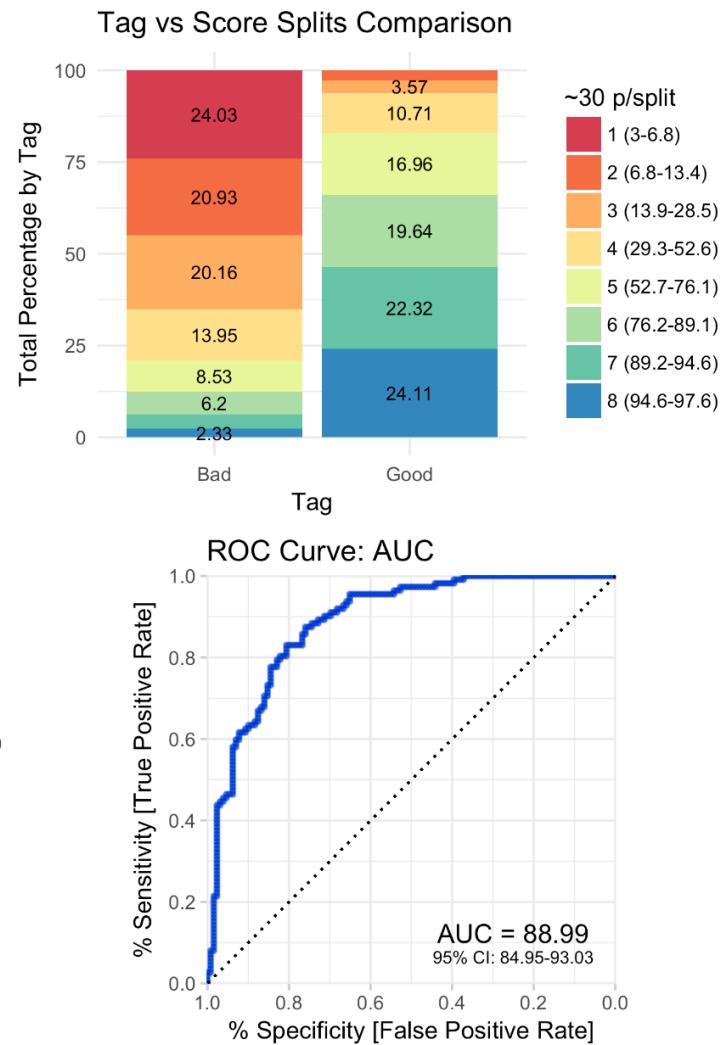
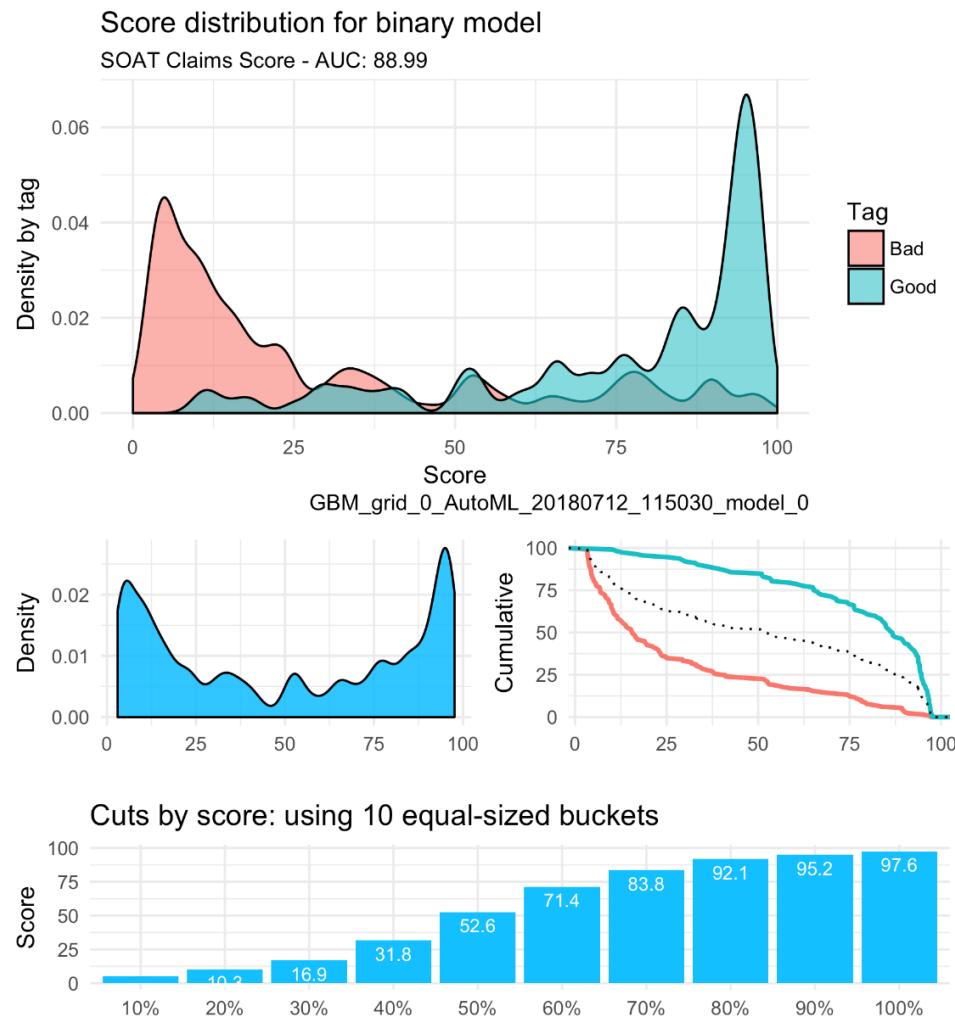
Where has deep learning visualization been used?

Application Domains & Models
A Vibrant Research Community

[Hohman, Fred, et al. 2018 "Visual analytics in deep learning: An interrogative survey for the next frontiers."]

MODEL PERFORMANCE

Performance: binary classifiers



[Source: <https://datascienceplus.com/machine-learning-results-one-plot-to-rule-them-all/>]

Confusion matrix

Compare predicted and actual class

Color code for more immediate interpretation



MODEL TRAINING

Tensorboard

- Default TF & Keras visualization tool
- Mostly used for managing and debugging training
- Different types of visualization including
 - ◆ **Graph dashboard:** Model graph, at conceptual and op-level
 - ◆ **Scalar dashboard:** tracking and visualizing loss and metrics, at batch and epoch level, during time
 - ◆ **Distributions and histograms dashboard:** Histograms of weights, biases, activations, gradients and how they change during training
 - ◆ **Tensorboard projector:** projecting embeddings on a lower dimensional space

Using tensorboard

- Tensorboard relies on log files that are created and updated during training
 - ◆ In Keras, this can be obtained using the `tf.keras.callbacks.Tensorboard` callback
- The Tensorboard application is launched separately from the training process
 - ◆ The interface can be opened using any web browser
 - ◆ Tensorboard can be stopped or launched independently of the training process
 - ◆ Integration with Colab is available

Conceptual graph

TensorBoard SCALARS GRAPHS TIME SERIES

Search nodes. Regexes supported.

Fit to Screen

Download PNG

Run (1) fit/20201229-221014/train

Tag (3) keras

Upload

Graph

Conceptual Graph

Profile

Trace inputs

Color Structure

Device

XLA Cluster

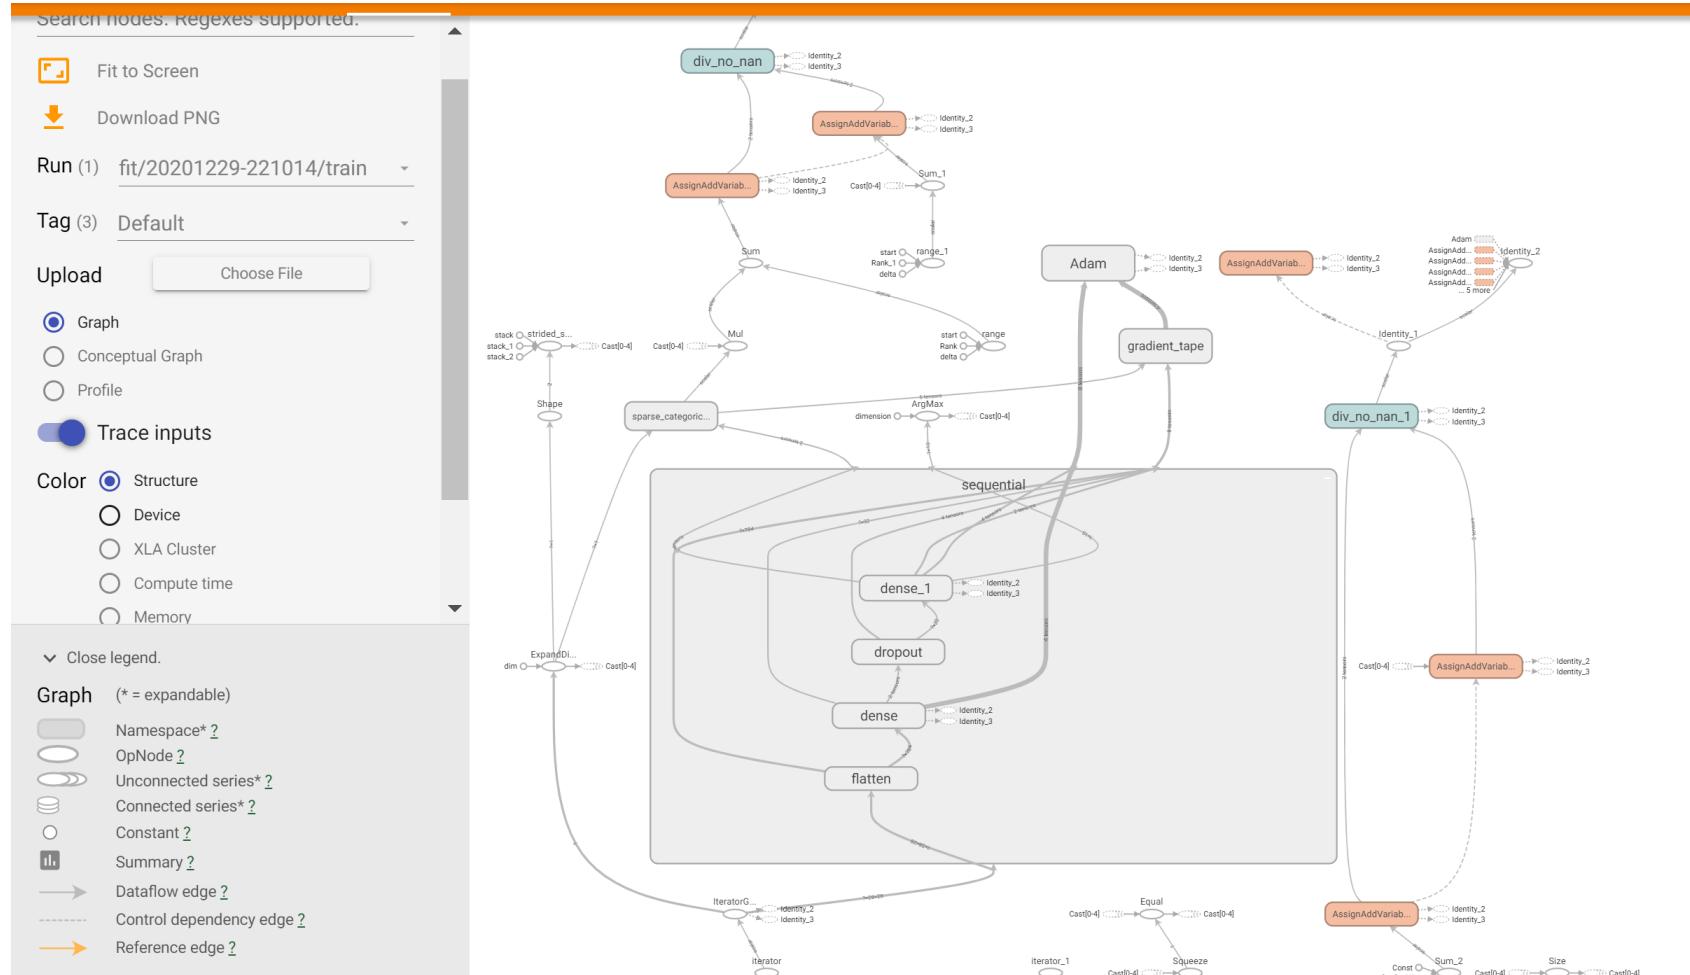
Compute time

Memory

sequential

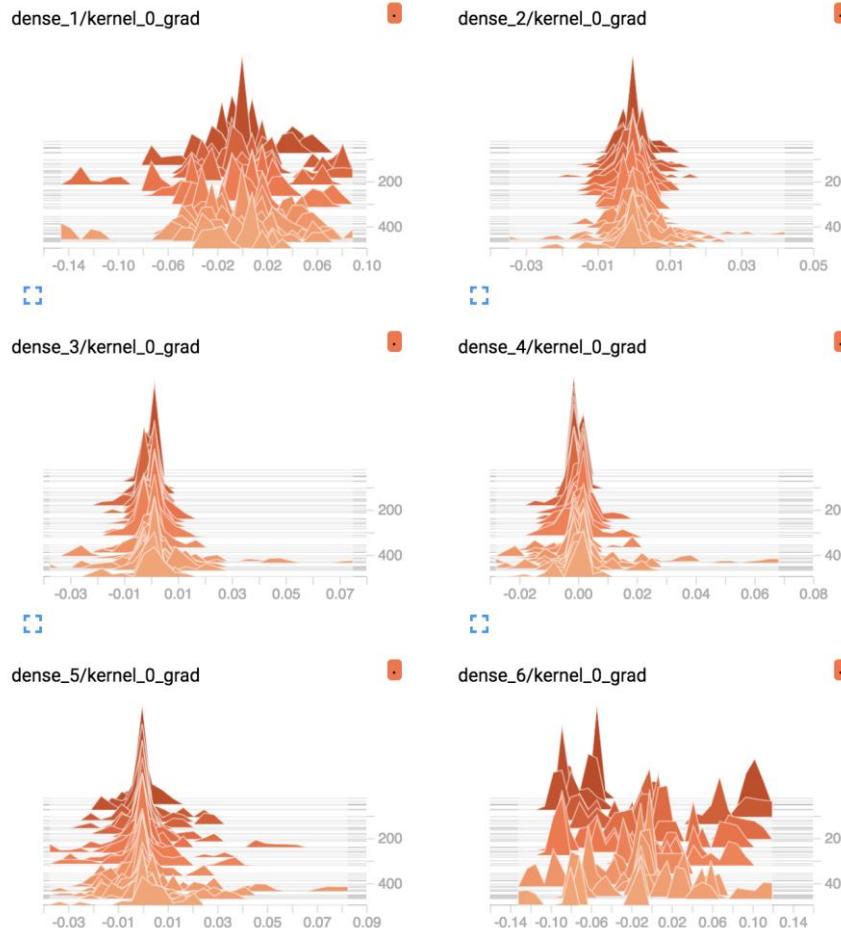
```
graph TD; flatten_in[flatten_in...] --> flatten[flatten]; flatten --> dense[dense]; dense --> dropout[dropout]; dropout --> dense1[dense_1];
```

Op-level graph

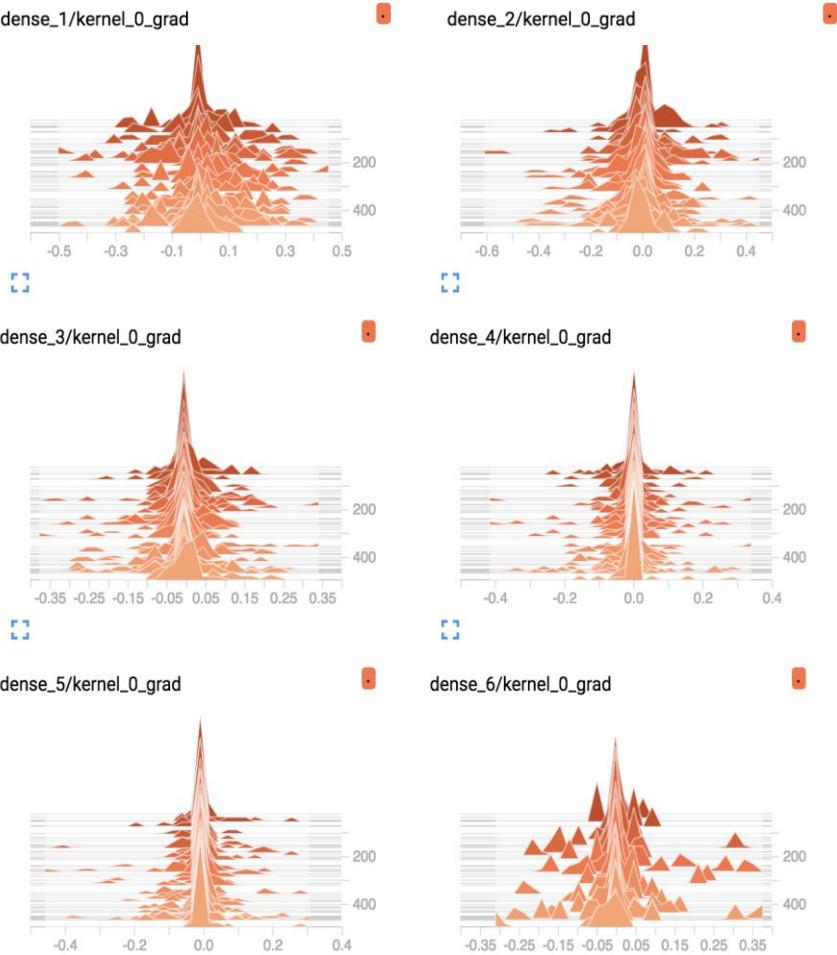


Distributions

MLP with tanh activations

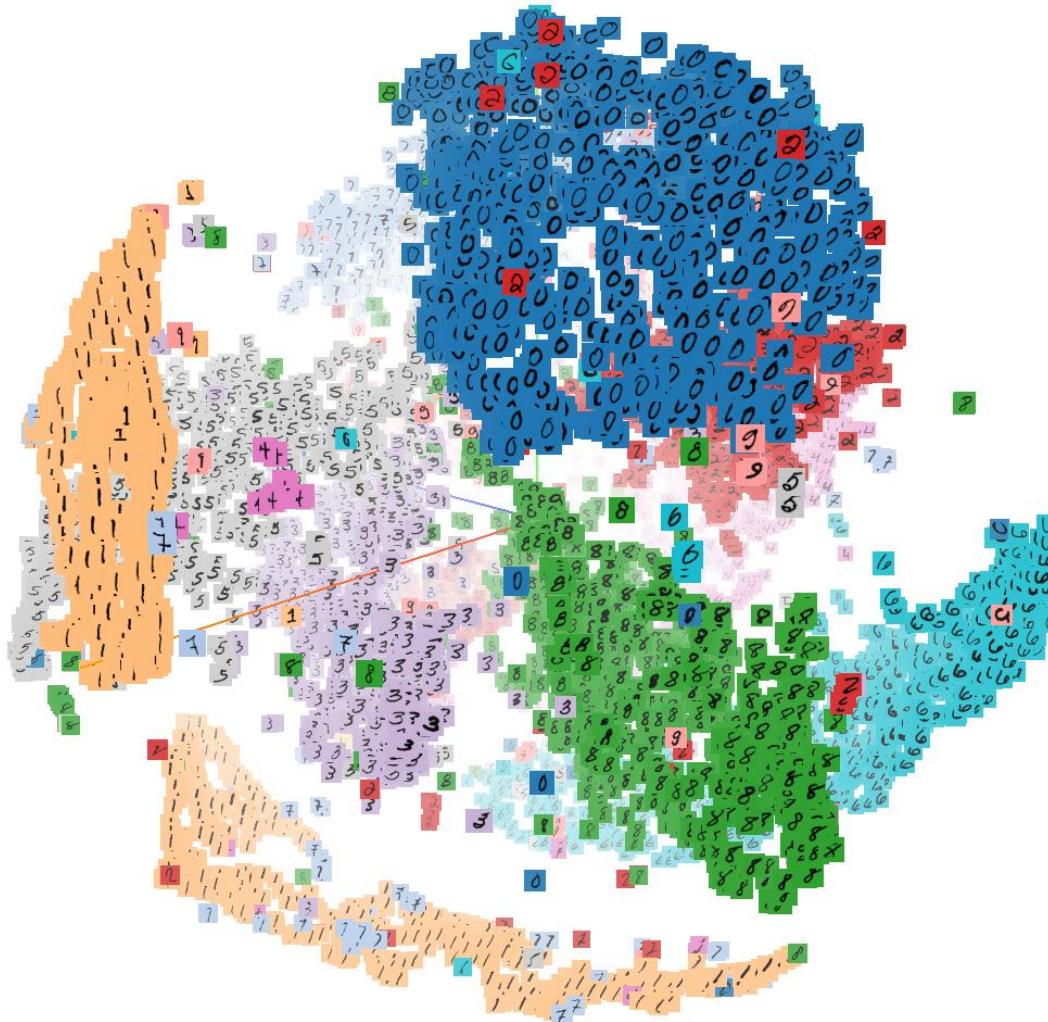


MLP with RELU activations



MODEL FEATURE SPACE

Embeddings

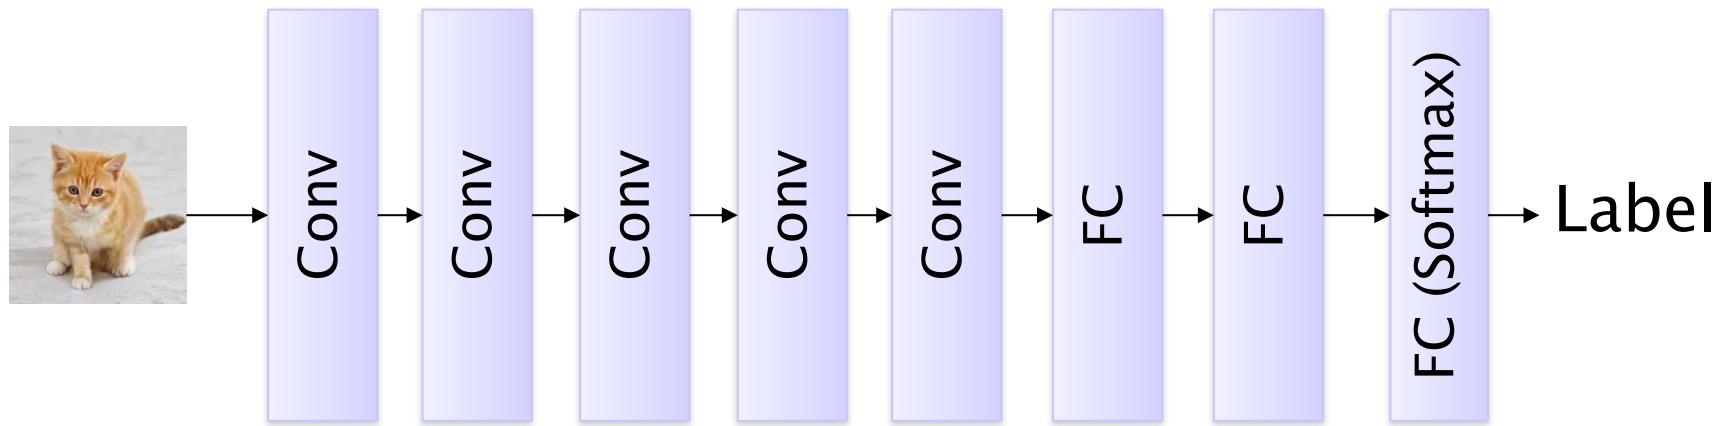


<https://projector.tensorflow.org/>

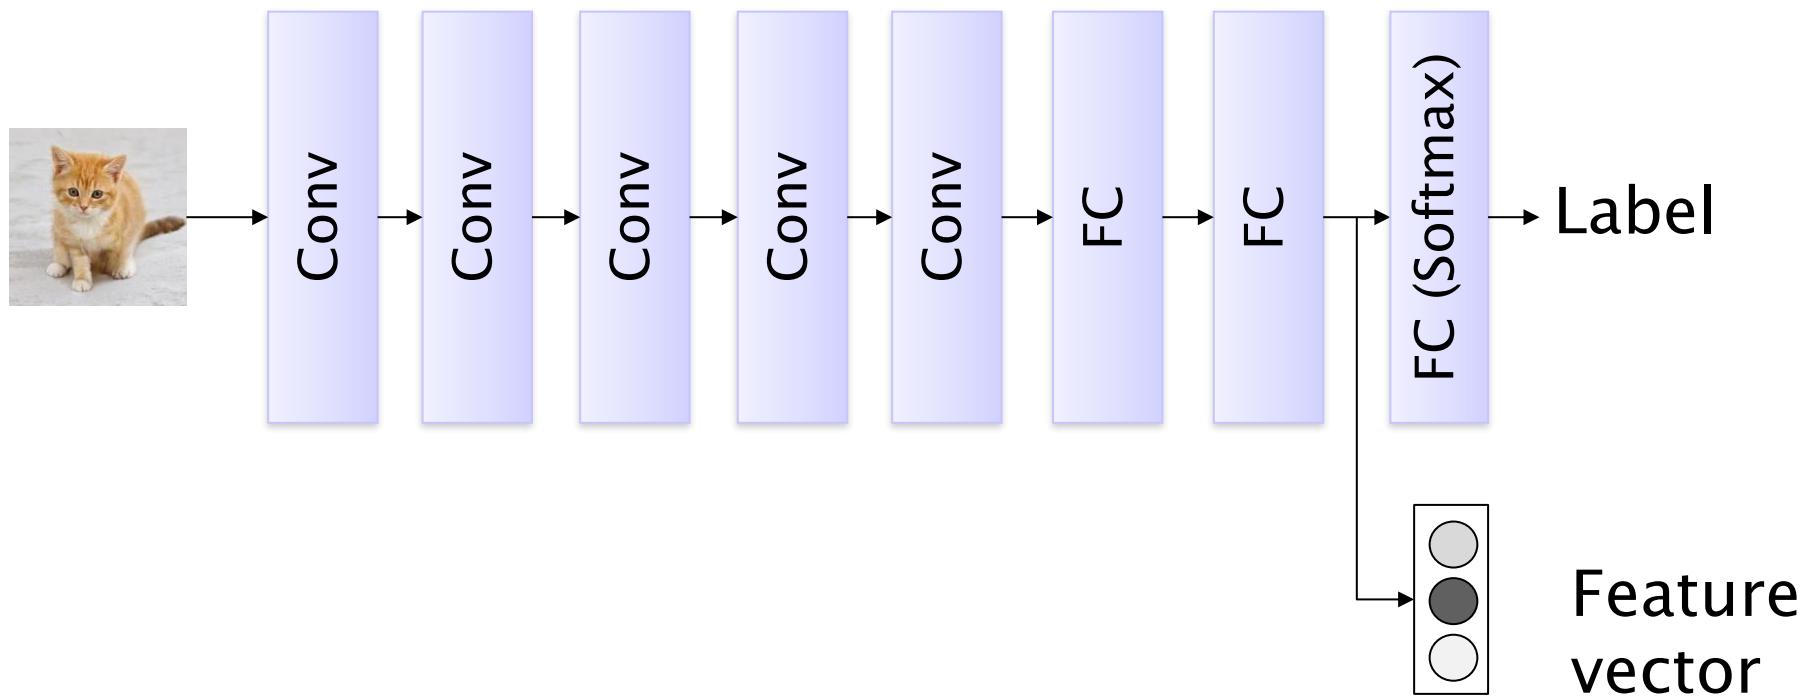
Visualizing embeddings

- Deep neural networks learn to map input data (e.g. images) into a lower-dimensional feature space
- Visualizing this feature space is important to understand its properties
 - ◆ Are samples from the same class nicely clustered?
 - ◆ Which examples are similar for the network?
- Dimensionality reduction techniques are needed to project the embedding space in 2 or 3 dimensions
- Available techniques include:
 - ◆ Principal Component Analysis (PCA)
 - ◆ t-Student Stochastic Neighbor Embeddings (tSNE)
 - ◆ Uniform Manifold Approximation and Projection (UMAP)

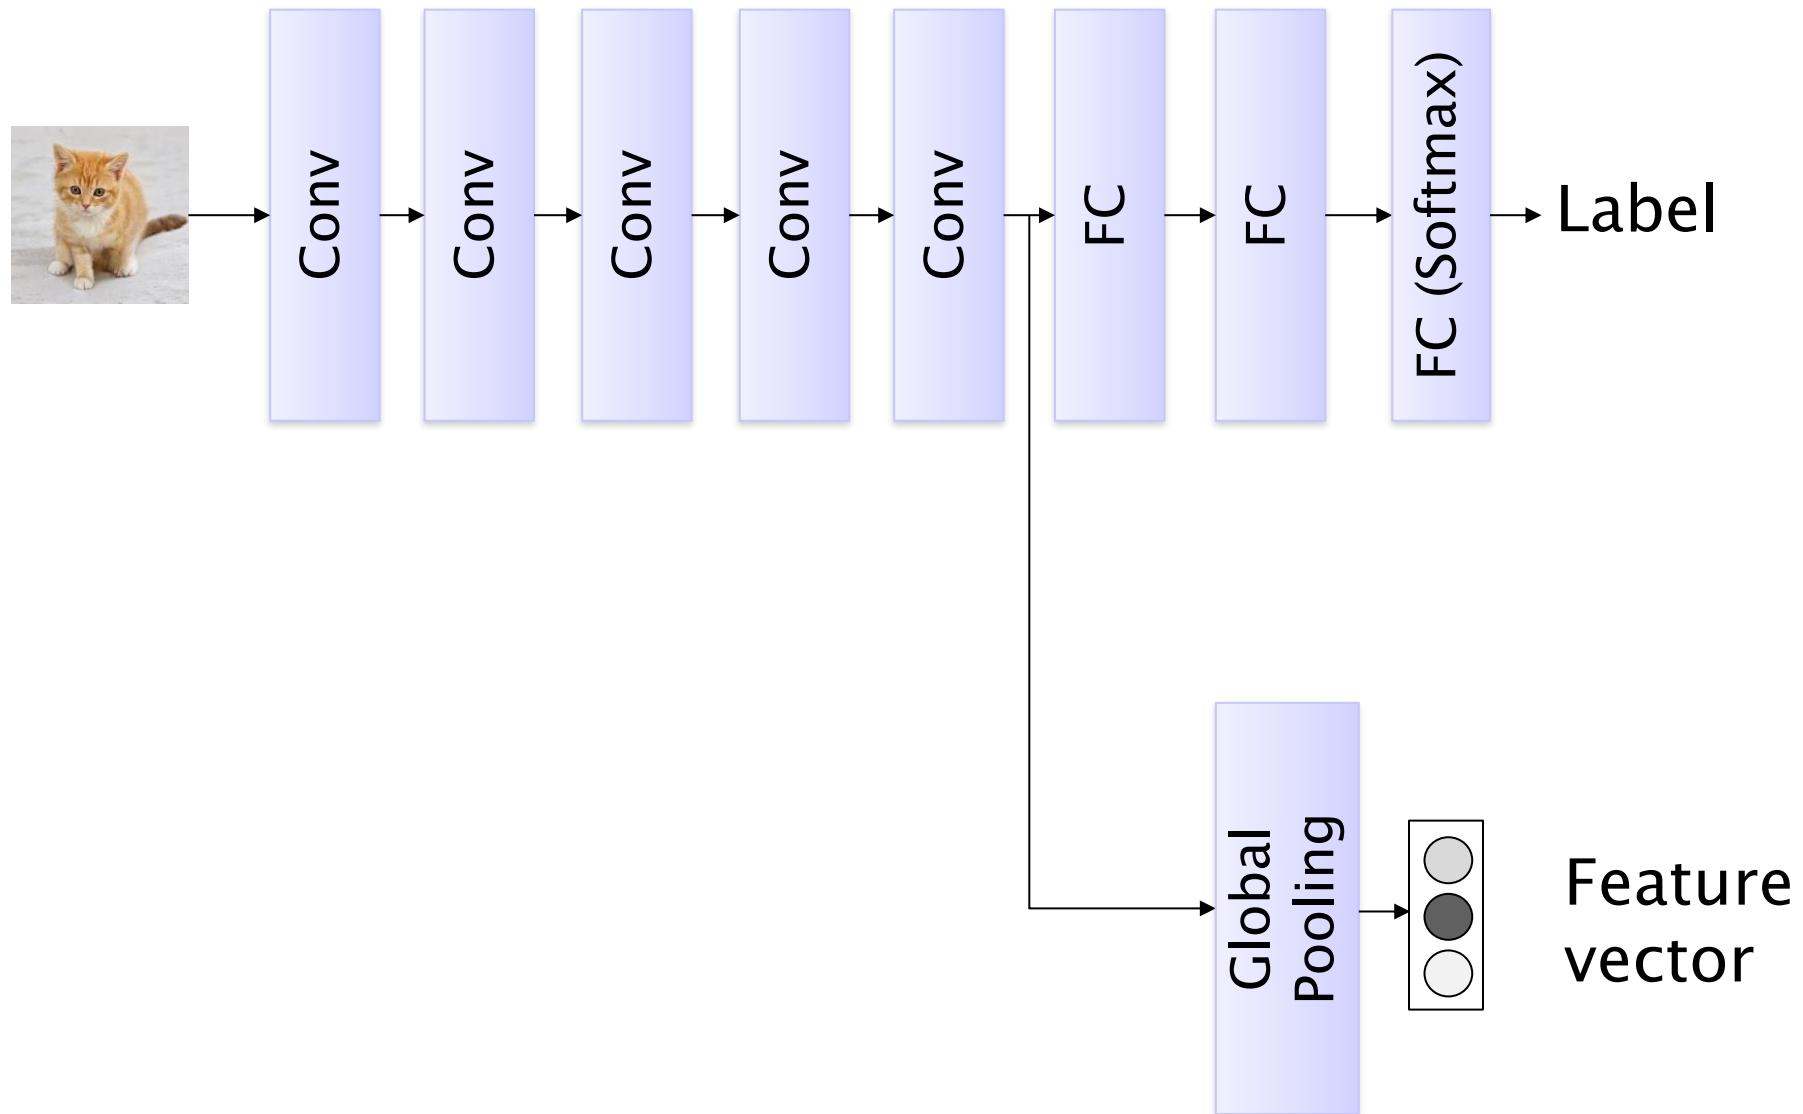
Extracting embeddings



Extracting embeddings



Extracting embeddings





tSNE embedding algorithm

1. (Optional) Use PCA to reduce the input dimensions into a smaller number
2. Construct a probability distribution over pairs of original high dimensional records: high probability \approx low distance
3. Define a similarity probability distribution of the points in the low-dimensional embedding
4. Minimize the distance between the two distributions using gradient descent method

tSNE embedding algorithm

- High-dim space:
 - ◆ Gaussian distribution
 - ◆ Focus on local neighbors
 - ◆ Depends on perplexity parameter (number of expected neighbours per point)
 - ◆ Perplexity typically between 5 and 50

$$p_{j|i} = \frac{\exp(-|x_i - x_j|/2\sigma_i^2)}{\sum_{k \neq i} \exp(-|x_i - x_k|/2\sigma_i^2)}$$

x_i, x_j, x_k represent the original samples (training set)

tSNE embedding algorithm

- High-dim space:
 - ◆ Gaussian distribution
 - ◆ Focus on local neighbors
 - ◆ Depends on perplexity parameter

$$p_{j|i} = \frac{\exp(-|x_i - x_j|/2\sigma_i^2)}{\sum_{k \neq i} \exp(-|x_i - x_k|/2\sigma_i^2)}$$

- Low-dim space:
 - ◆ Student-t distribution
 - ◆ Heavy tails compensates for lower dimensionality

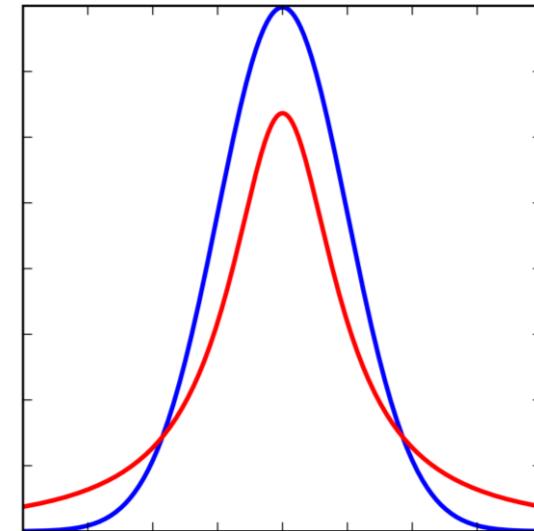
y_i, y_j, y_k represent the projections on low-dim space (output)

$$q_{ij} = \frac{f(|y_i - y_j|)}{\sum_{k \neq i} f(|y_i - y_k|)} \text{ with } f(z) = \frac{1}{1+z^2}$$

tSNE embedding algorithm

- High-dim space:
 - ◆ Gaussian distribution
 - ◆ Focus on local neighbors
 - ◆ Depends on perplexity parameter

$$p_{j|i} = \frac{\exp(-|x_i - x_j|^2 / 2\sigma_i^2)}{\sum_{k \neq i} \exp(-|x_i - x_k|^2 / 2\sigma_i^2)}$$



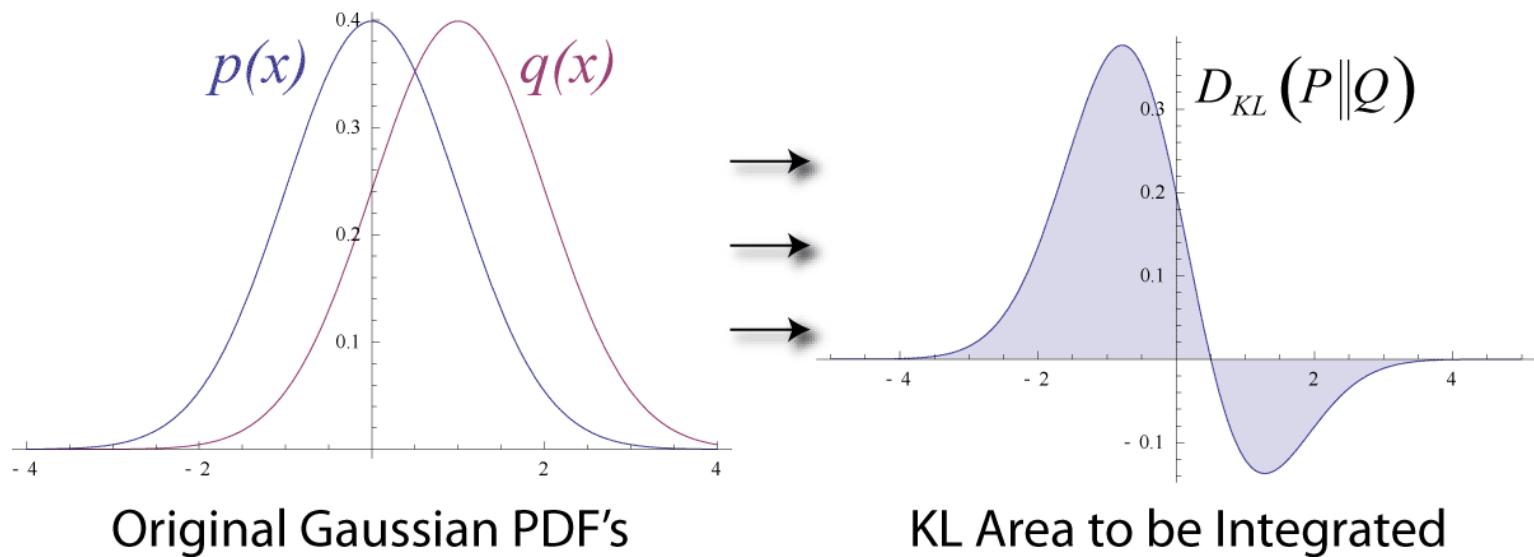
- Low-dim space:
 - ◆ Student-t distribution
 - ◆ Heavy tails compensates for lower dimensionality

$$q_{ij} = \frac{f(|y_i - y_j|)}{\sum_{k \neq i} f(|y_i - y_k|)} \text{ with } f(z) = \frac{1}{1+z^2}$$

Kullback–Leiber divergence

Measures how one probability distribution P differs from a second, reference probability distribution Q

$$KL(P \parallel Q) = \int P(x) \log \frac{P(x)}{Q(x)} dx$$



tSNE embedding algorithm

- Minimize Kullback–Leiber divergence by gradient descent optimization

$$KL(P||Q) = \sum_{i,j} p_{ij} \log \frac{p_{ij}}{q_{ij}}$$

- KL divergence is not symmetrical
 - ◆ large p (“similar” points in input space) modelled with small q (“dissimilar” points in output space) → big penalty
 - ◆ small p (“dissimilar” points in input space) modelled with large q (“similar” points in output space) → small penalty
 - ◆ preserves local (not global) structure: distance between clusters may not be meaningful

tSNE hyperparameters

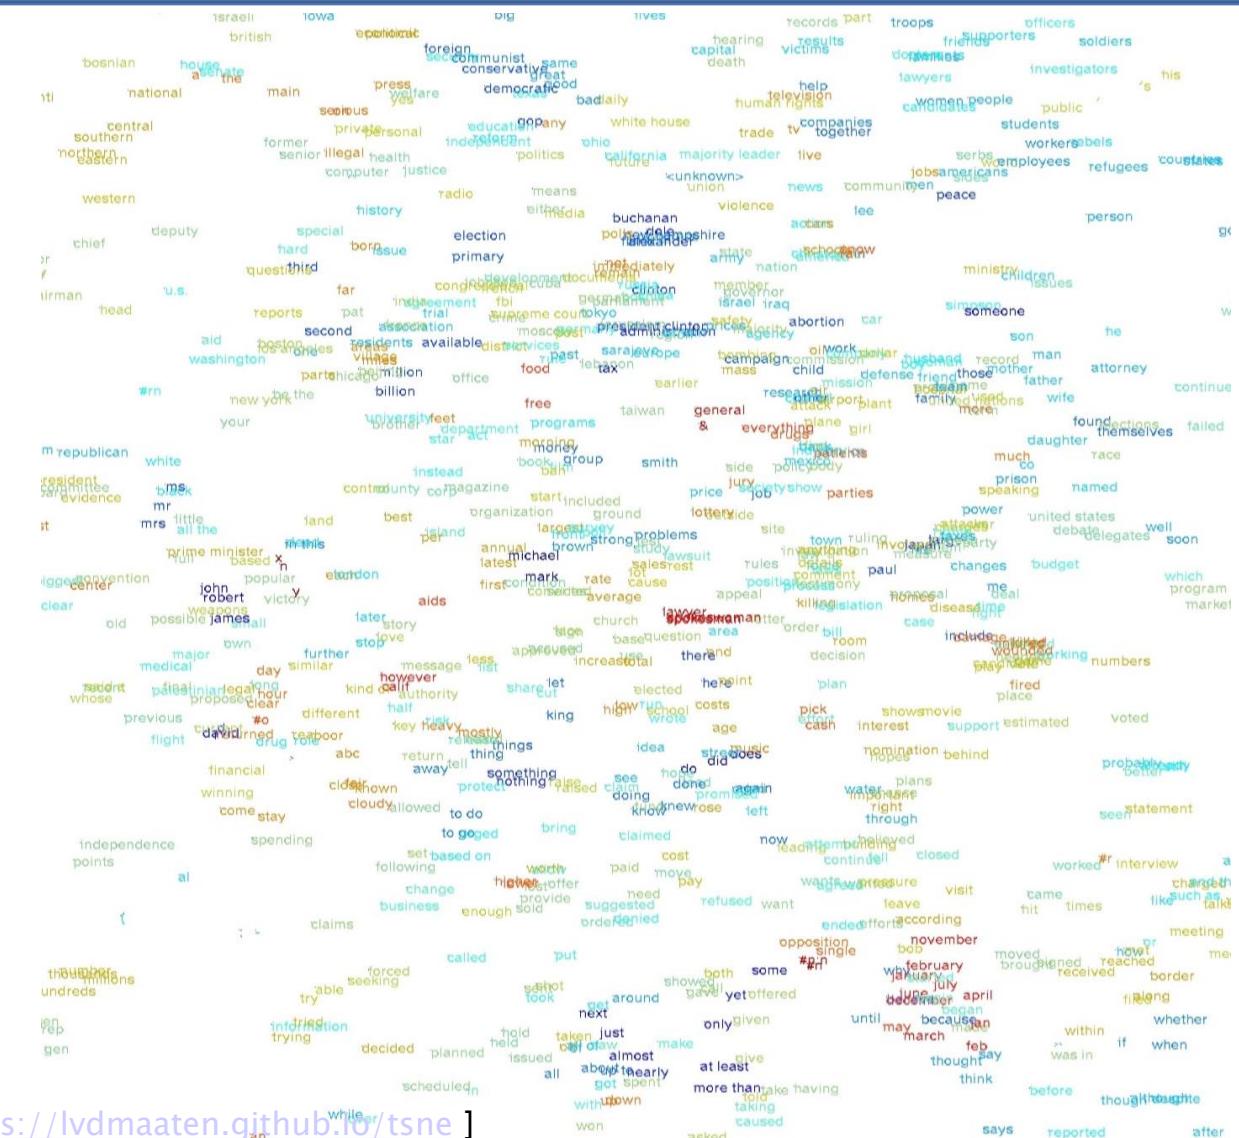
- Perplexity:
 - ◆ balance attention between local and global aspects of the dataset
 - ◆ best guess about the number of close neighbors
 - ◆ good values between 5 and 50, but important to try different values in real parameters
 - ◆ must be lower than the number of input records
- Learning rate, training iterations:
 - ◆ train until convergence
 - ◆ beware that multiple iterations will results in different visualizations

tSNE example: Caltech101



[source: <https://lvdmaaten.github.io/tsne>]

tSNE example: word embeddings



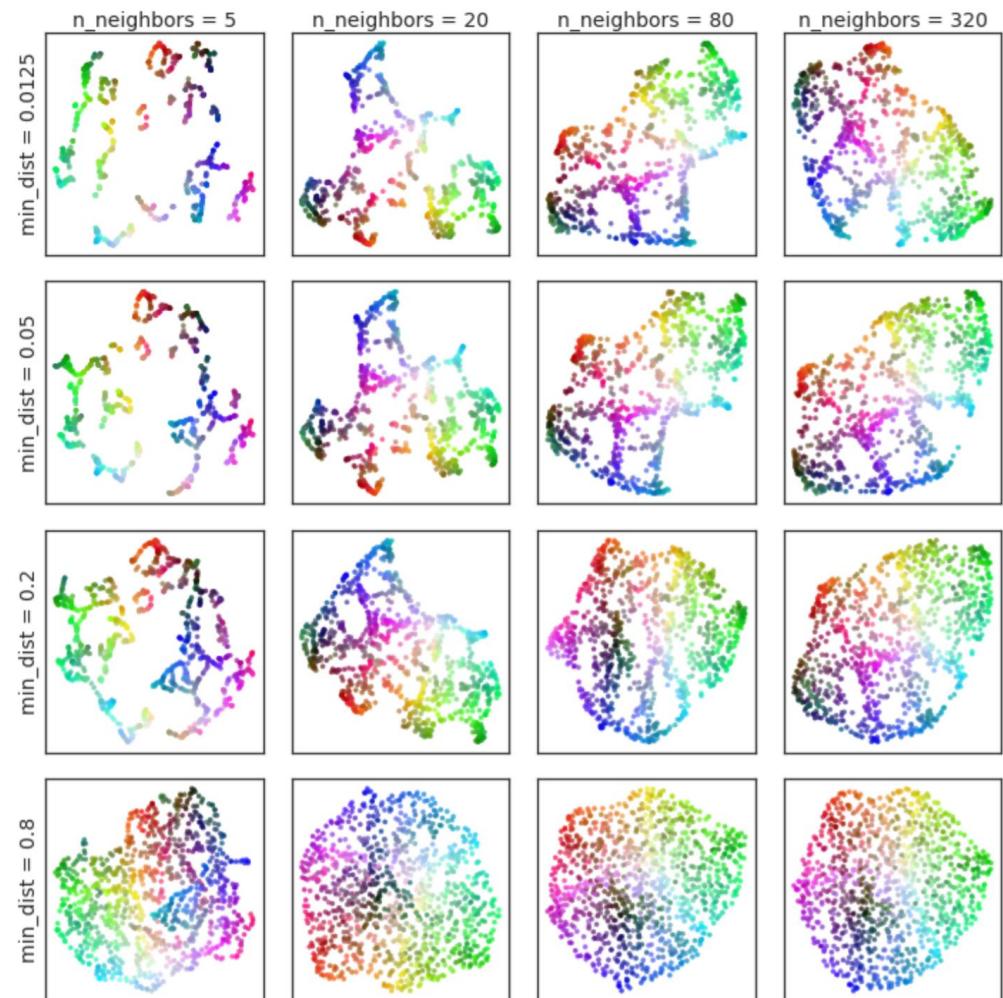
[source: <https://lvdmaaten.github.io/tsne/>]

Interpreting embeddings

- Manifold learning algorithms assume that there is a structure in the data
- Highly non-linear embeddings can fit a structure to the noise of a dataset – similarly to humans!
- Care must be taken with small sample sizes of noisy data, or data with only large scale manifold structure
- Shape may not be always preserved.
 - ◆ See <https://distill.pub/2016/misread-tsne/>

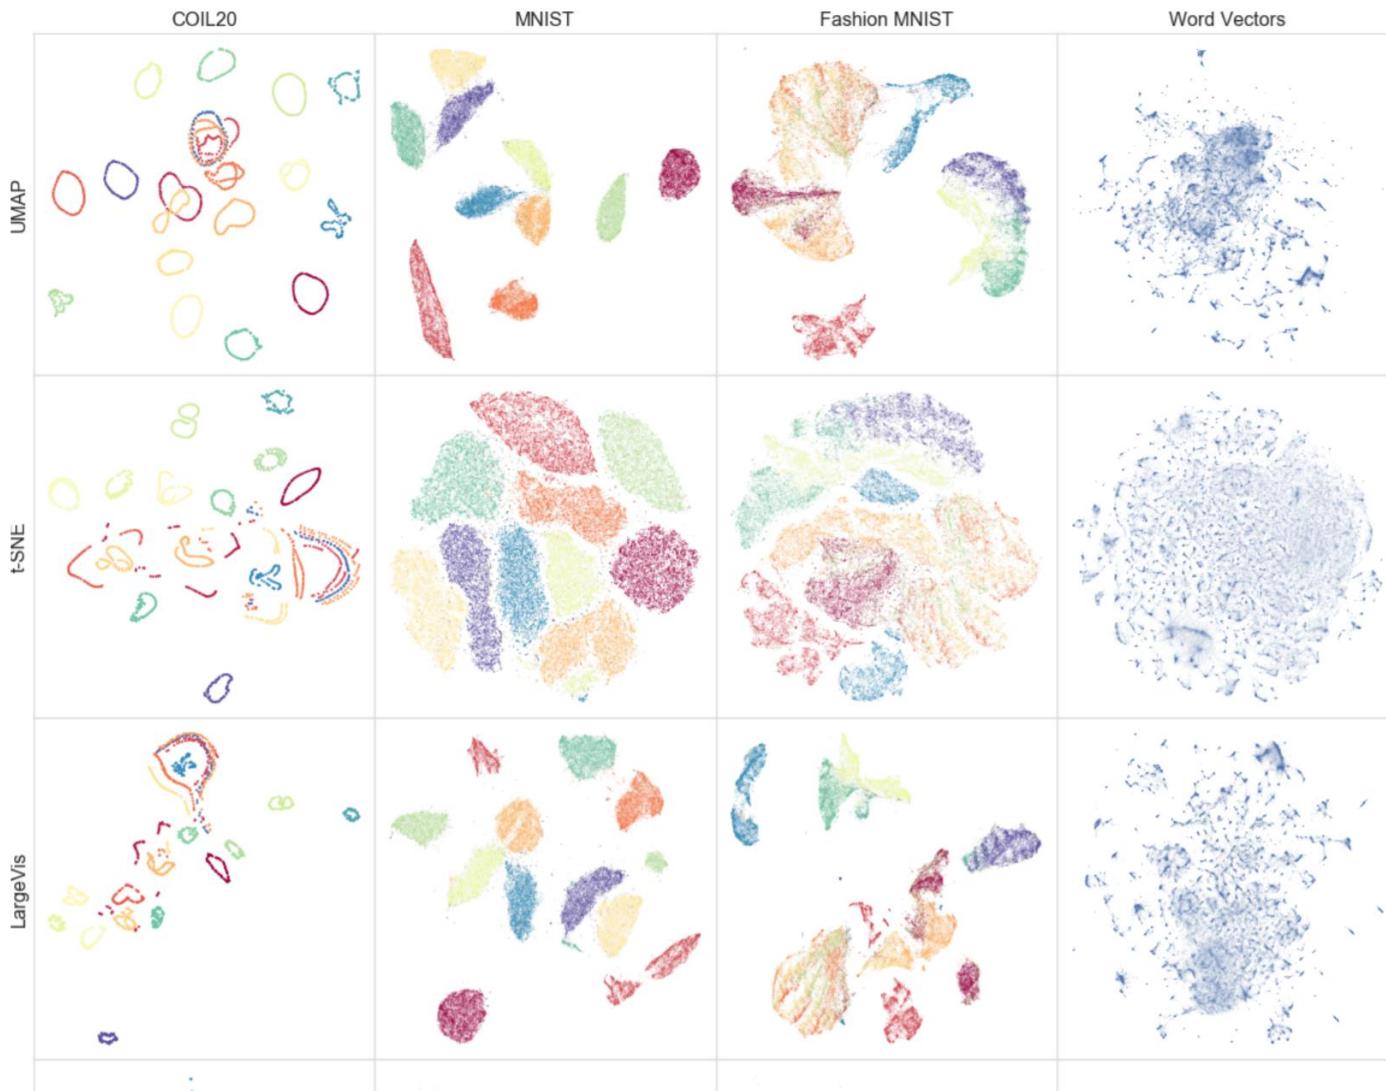
Interpreting embeddings

- Example: randomly sampled data in the RGB domain (3 to 2 dimensions)
- UMAP trained with different choices of hyper-parameters shows emerging spurious structures



[McInnes et al. 2018. Umap: Uniform manifold approximation and projection for dimension reduction. *arXiv preprint arXiv:1802.03426*.]

Interpreting embeddings (II)

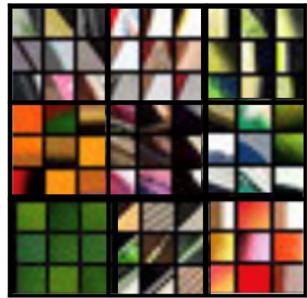


FEATURE VISUALIZATION

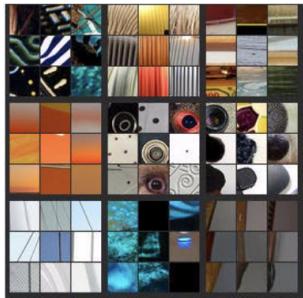
Visualizing features

- Visualizing the activations of individual neurons may be long and not very informative (networks are large and deep, activation maps difficult to interpret and sparse)
- There are fundamentally two ways of visualizing individual neurons
 - ◆ Finding the patterns that optimize its response
 - ◆ Finding images (or patches) that activate (high-level) neurons: can we see any pattern emerge?
- Visualizing features before/after training shows “repurposing” of individual neurons

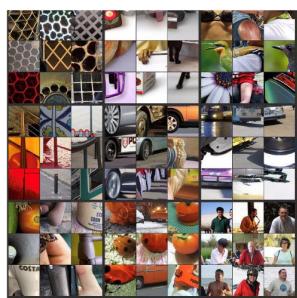
Visualize what network is learning



Layer 1



Layer 2



Layer 3



Layer 4



Layer 5

[Zeiler and Fergus., 2013, Visualizing and understanding convolutional networks]

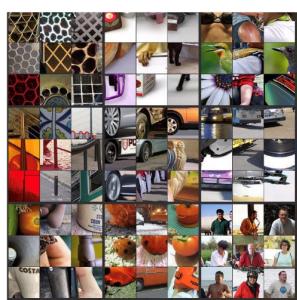
Visualize what network is learning



Layer 1



Layer 2



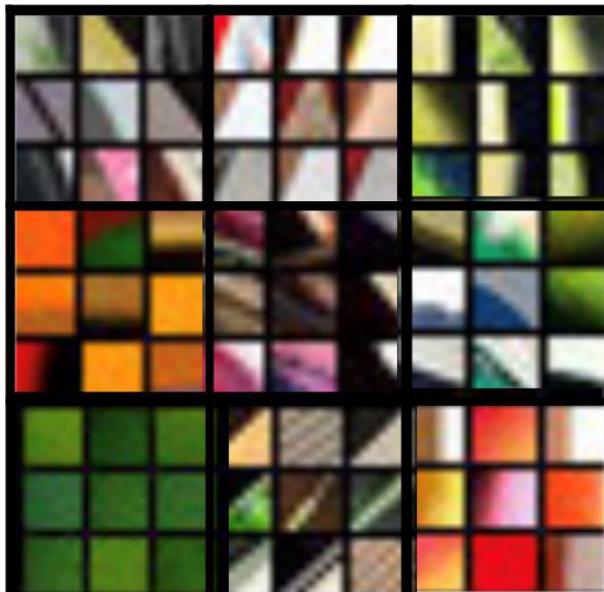
Layer 3



Layer 4



Layer 5



[Zeiler and Fergus., 2013, Visualizing and understanding convolutional networks]

Visualize what network is learning



Layer 1



Layer 2



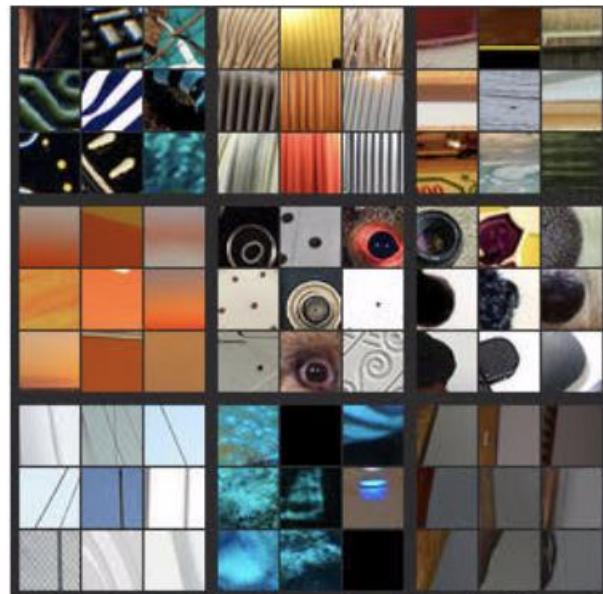
Layer 3



Layer 4



Layer 5



[Zeiler and Fergus., 2013, Visualizing and understanding convolutional networks]

Visualize what network is learning



Layer 1



Layer 2



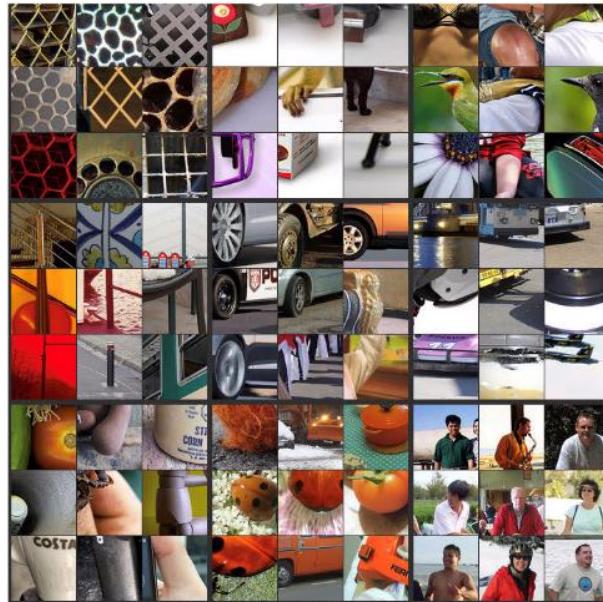
Layer 3



Layer 4

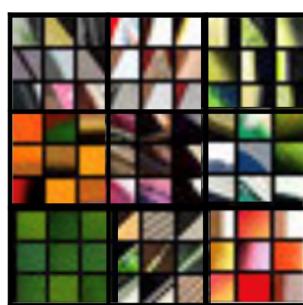


Layer 5

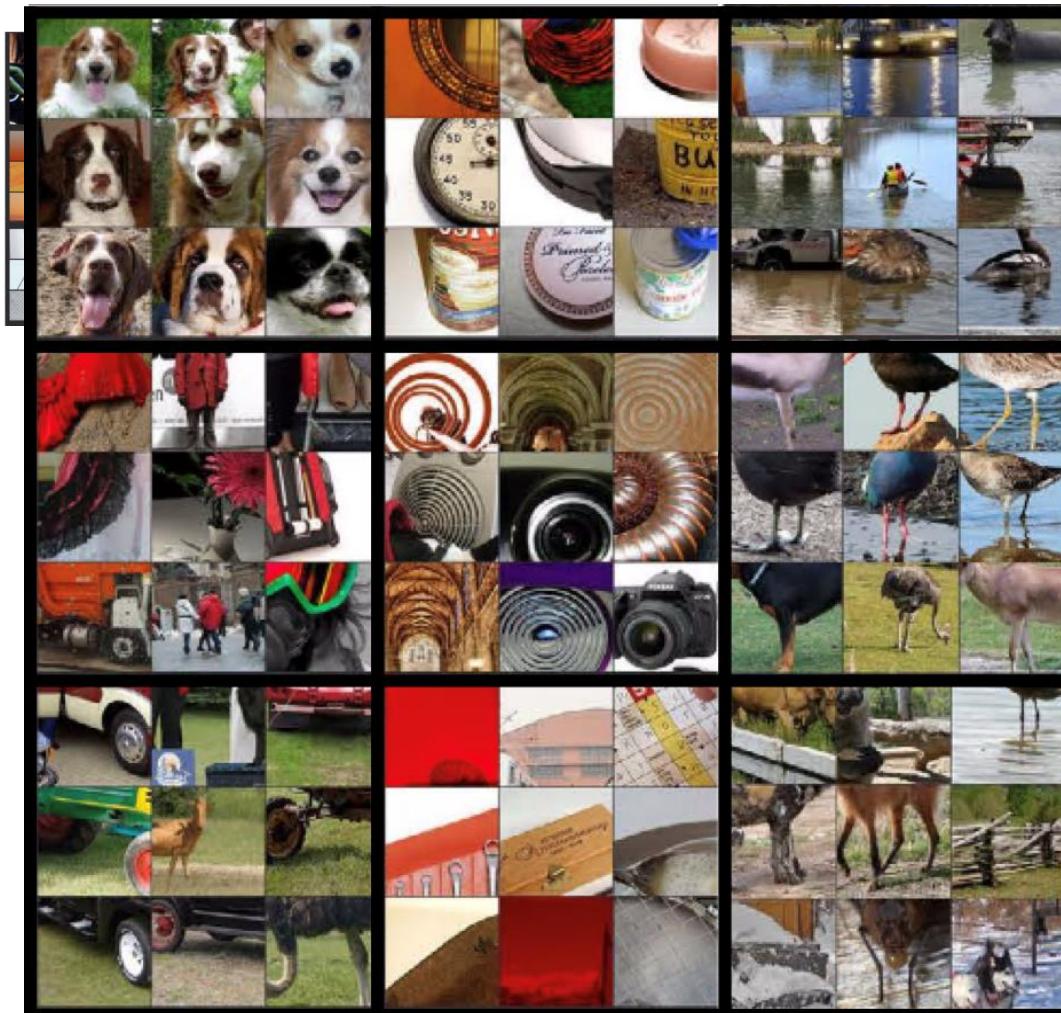


[Zeiler and Fergus., 2013, Visualizing and understanding convolutional networks]

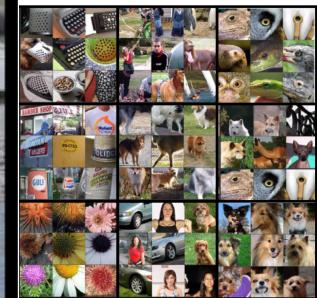
Visualize what network is learning



Layer 1

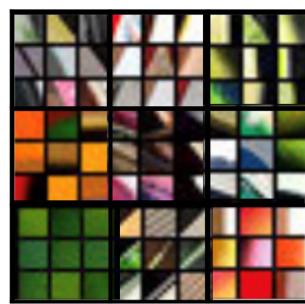


Layer 5

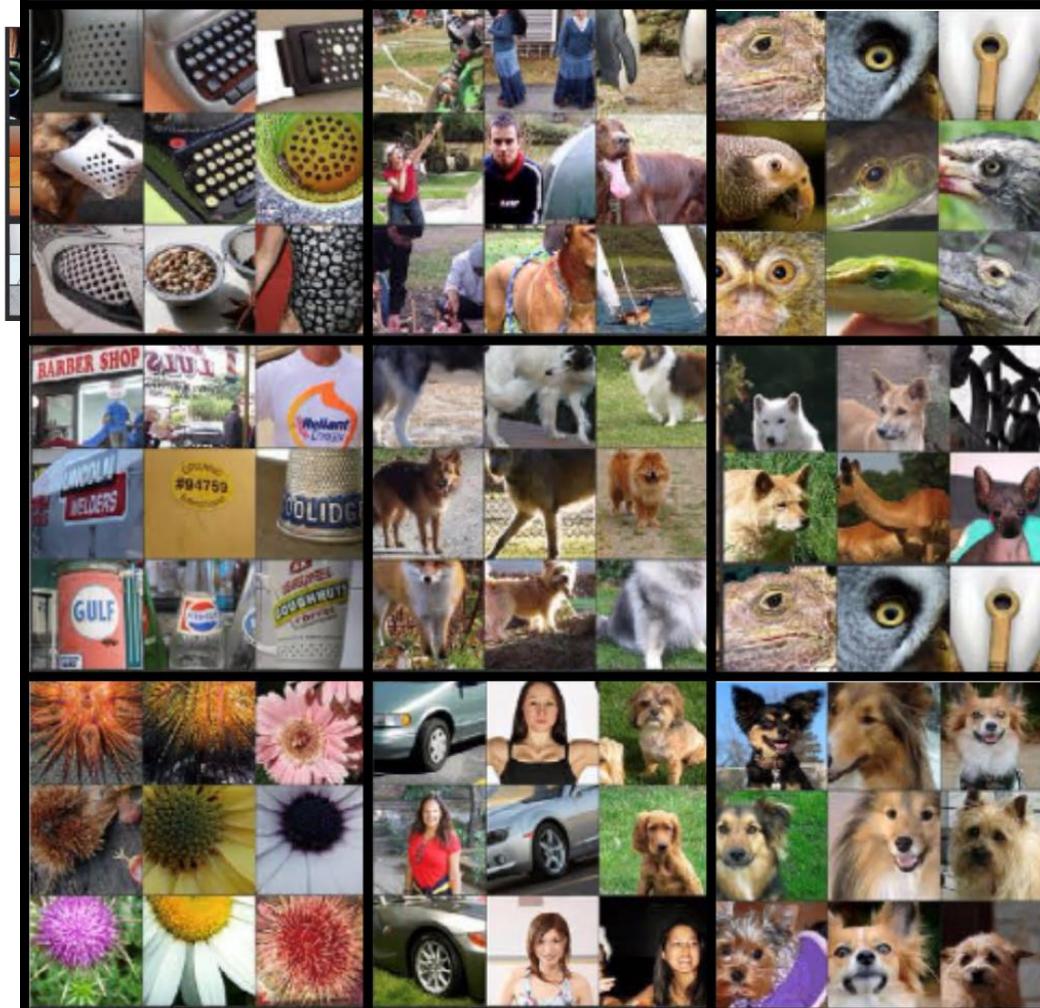


[Zeiler and Fergus., 2013, Visualizing and understanding convolutional networks]

Visualize what network is learning



Layer 1

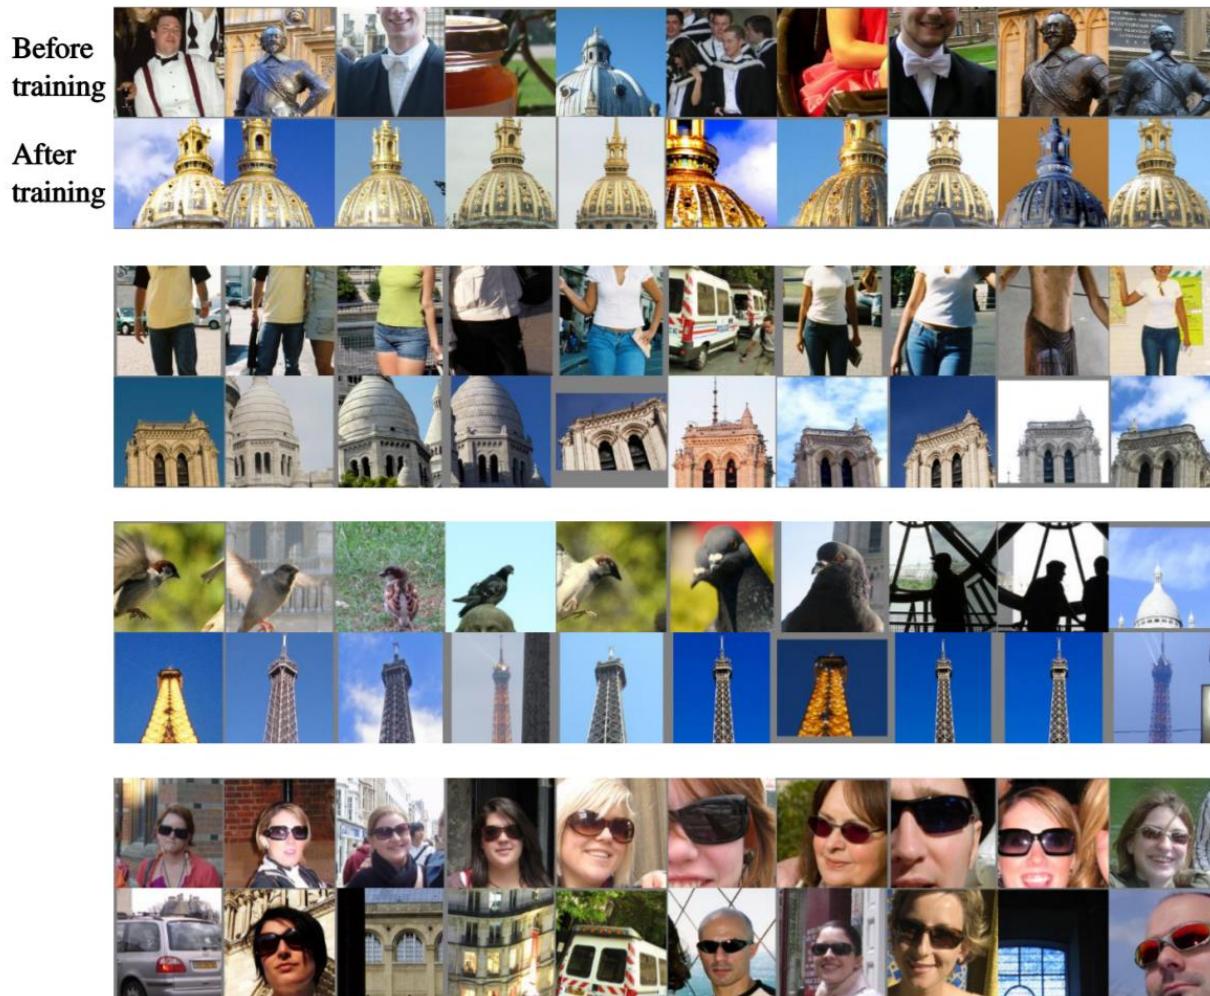


Layer 5



[Zeiler and Fergus., 2013, Visualizing and understanding convolutional networks]

Re-purposing



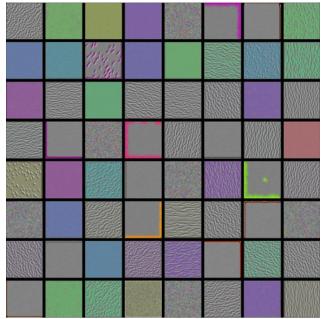
[Gordo, 2016, Deep Image Retrieval: Learning global representations for image search]

Visualization by gradient ascent

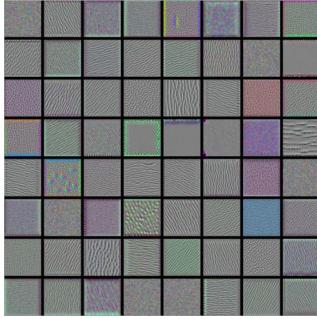
- Another easy way to inspect the filters learned by convolutional networks is to display the visual pattern that each filter is meant to respond to.
- This can be done with **gradient ascent** in input space: applying gradient descent to the value of the input image to maximize the response of a specific filter, starting from a random input image.
- The resulting input image will be one that the chosen filter is maximally responsive to.



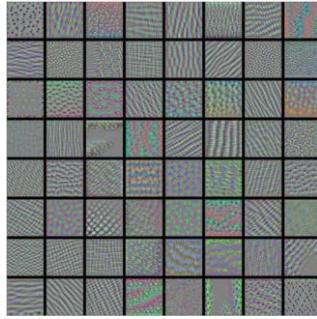
Gradient ascent



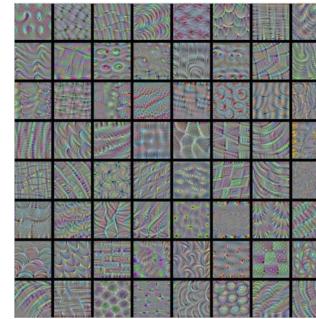
Layer 1



Layer 2



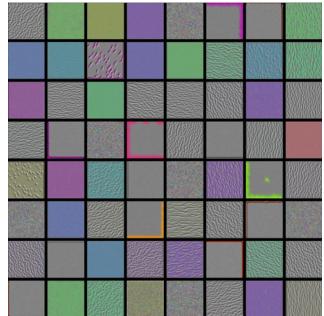
Layer 3



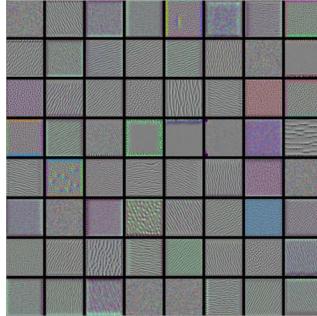
Layer 4



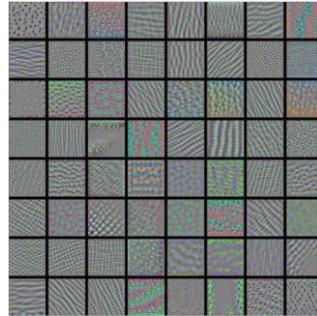
Gradient ascent



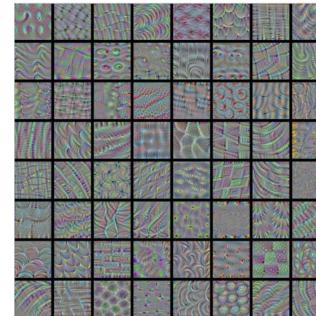
Layer 1



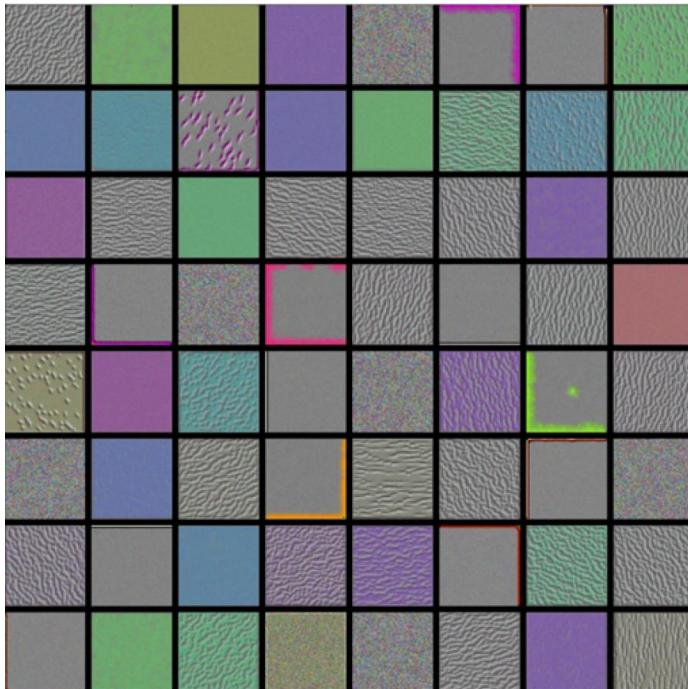
Layer 2



Layer 3

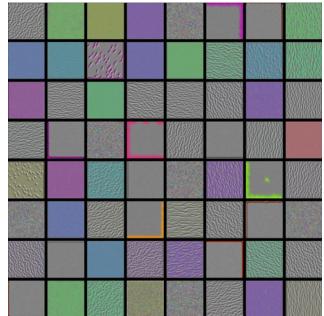


Layer 4

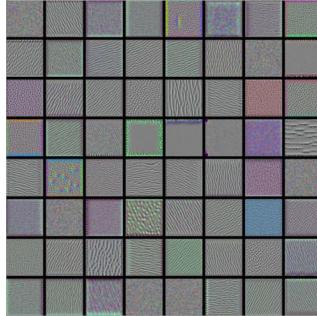




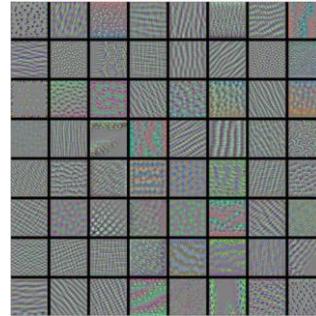
Gradient ascent



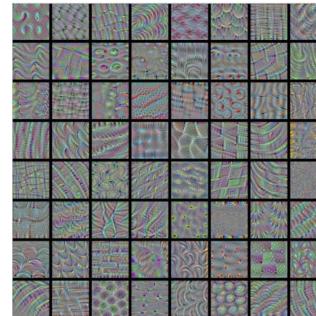
Layer 1



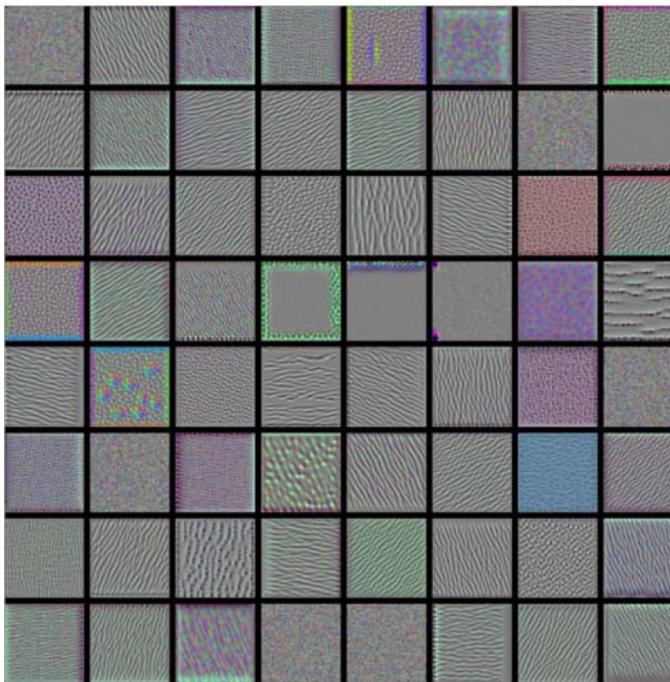
Layer 2



Layer 3

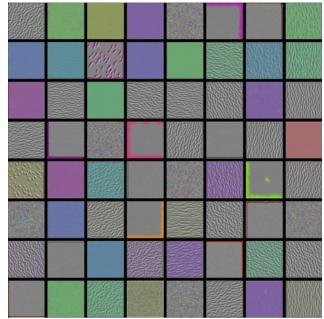


Layer 4

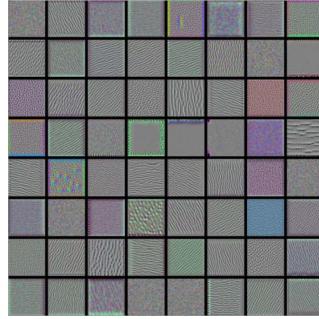




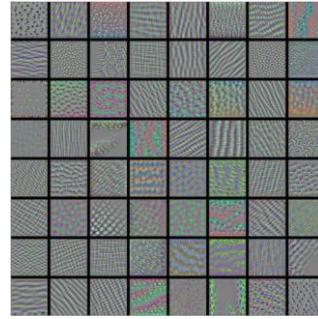
Gradient ascent



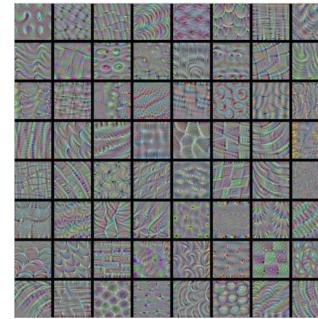
Layer 1



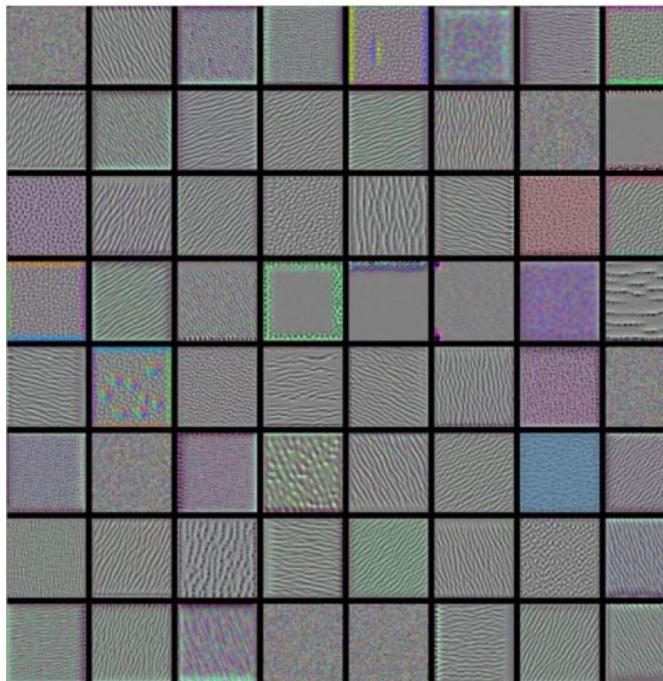
Layer 2



Layer 3

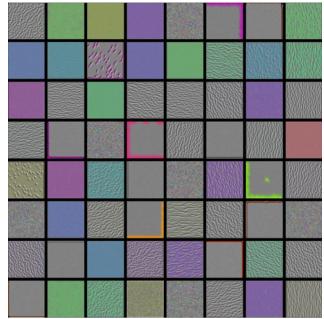


Layer 4

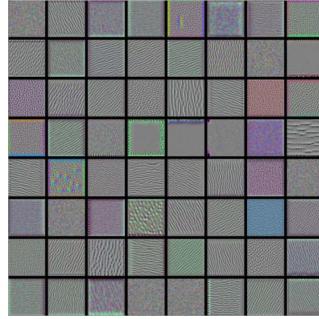




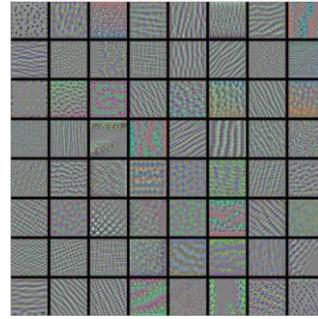
Gradient ascent



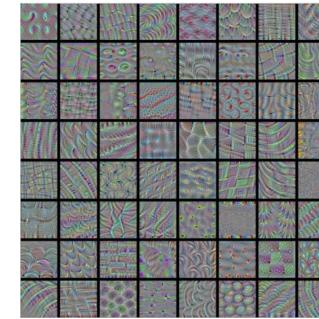
Layer 1



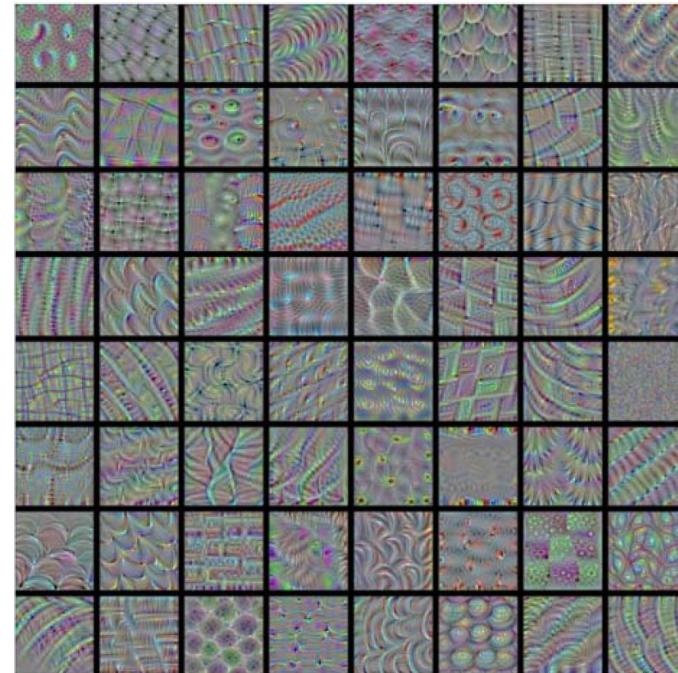
Layer 2



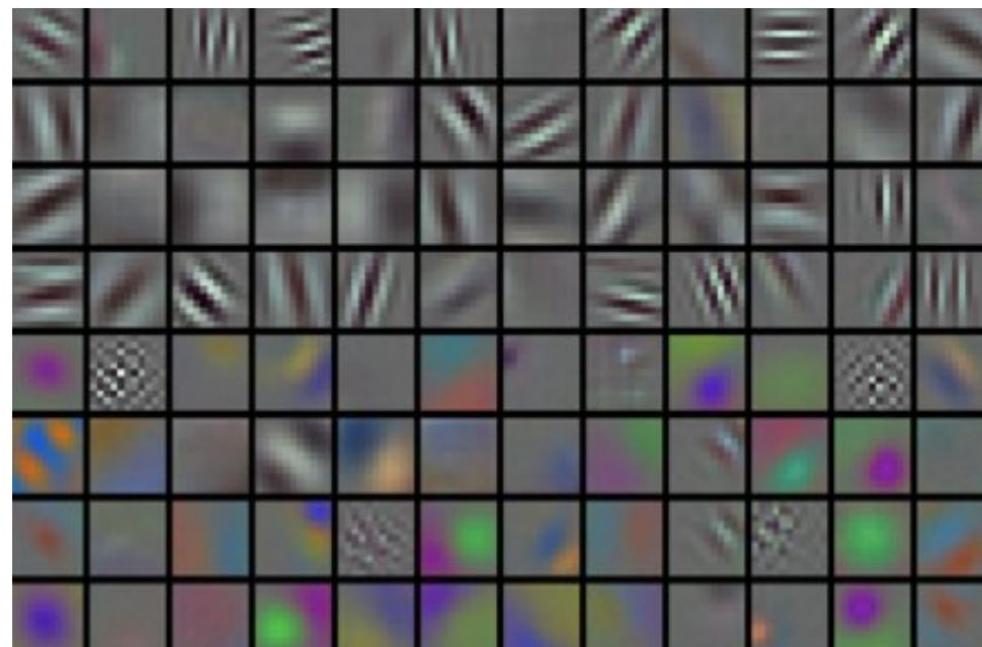
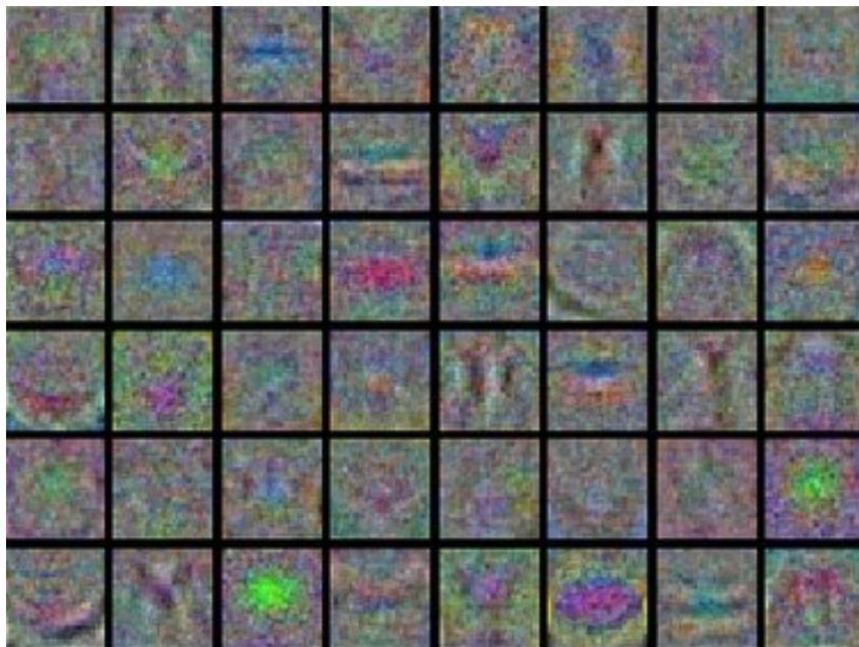
Layer 3



Layer 4



Visualize filter kernels



Filter kernels can be visualized directly as images
Noisy features (left) are usually the sign of poor training
(low learning rate, low number of epochs)

[Source: <https://cs231n.github.io/neural-networks-3/#vis>]

MODEL EXPLANATION

Model explanation

- “Explaining a prediction” → presenting a **qualitative** and possibly **quantitative** understanding of the **relationship** between the instance’s components (e.g. words in text, patches in an image) and the model’s prediction.
- Explanations are important to determine
 - ◆ Why and how a model is failing
 - ◆ Whether a model is correct “for the right reasons”
- Most techniques are relevance-based: which features activate a certain neuron/output? which features are most relevant for a certain classification?

Model explanation (III)

- Visualization is a key techniques for enhancing interpretability especially for CNNs
- Most techniques identify pixels which are mostly relevant to a given classification (“where” the network is looking)
- Different types of visualization techniques have been proposed
 - ◆ gradient-based
 - ◆ attribution-based
 - ◆ localization-based
 - ◆ perturbation-based methods

Gradient-based approaches

- Find the derivatives of class scores with respect to pixel intensities
 - ◆ Deconvolution
 - ◆ Guided Backpropagation
 - ◆ SmoothGrad
 - ◆ etc..



Attribution-based approaches



- For each pixel, measure its relevance to the score
 - ◆ LRP, DeepTaylor, PatternAttribution, LIME
 - The key idea is to decompose the output in terms of contributions from its inputs
 - Some techniques (e.g., LIME) are applicable to other types of input features other than images
-



GRAD-CAM

- GRAD-CAM computes a heatmap which assess the importance of each pixel with respect to a given class c , with output score y^c
- Let A^k be the activation of the k^{th} neuron in a convolutional layer
- First, the importance of each neuron is computed

$$\alpha_c^k = \frac{1}{Z} \sum_i \sum_j \frac{\partial y^c}{\partial A_{ij}^k}$$



GRAD-CAM

- GRAD-CAM computes a heatmap which assess the importance of each pixel with respect to a given class c , with output score y^c
- Let A^k be the activation of the k^{th} neuron in a convolutional layer
- First, the importance of each neuron is computed

Global Average
Pooling

$$\alpha_c^k = \frac{1}{Z} \sum_i \sum_j \frac{\partial y^c}{\partial A_{ij}^k}$$



GRAD-CAM

- GRAD-CAM computes a heatmap which assess the importance of each pixel with respect to a given class c , with output score y^c
- Let A^k be the activation of the k^{th} neuron in a convolutional layer
- First, the importance of each neuron is computed

Global Average Pooling $\alpha_c^k = \frac{1}{Z} \sum_i \sum_j \frac{\partial y^c}{\partial A_{ij}^k}$ Gradients of y^c w.r.t to filter k



GRAD-CAM

- GRAD-CAM computes a heatmap which assess the importance of each pixel with respect to a given class c , with output score y^c
- Let A^k be the activation of the k^{th} neuron in a convolutional layer
- First, the importance of each neuron is computed

Global Average Pooling $\alpha_c^k = \frac{1}{Z} \sum_i \sum_j \frac{\partial y^c}{\partial A_{ij}^k}$ **Gradients of y^c w.r.t to filter k**

- The heatmap weights activations by their importance

$$L^c = \text{ReLU} \left(\sum_k \alpha_c^k A^k \right)$$

GRAD-CAM Example (I)

Grad-CAM for "Cat"



Grad-CAM for "Dog"



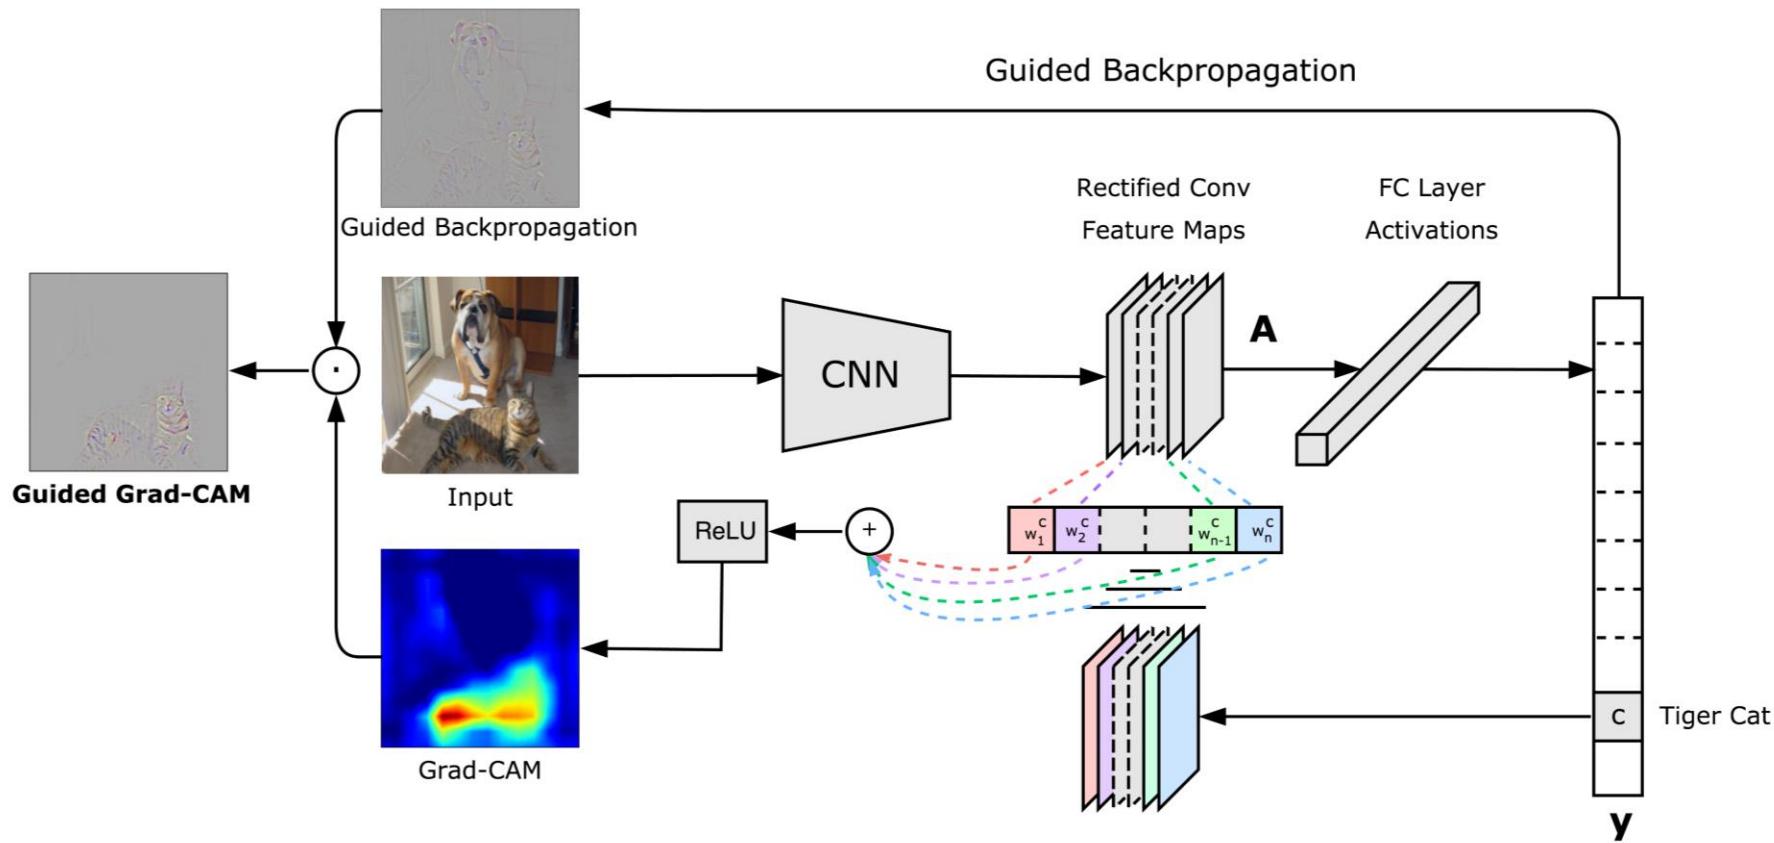
GRAD-CAM Example (II)



Visualization may identify biases in the trained model, then corrected by revising the training set

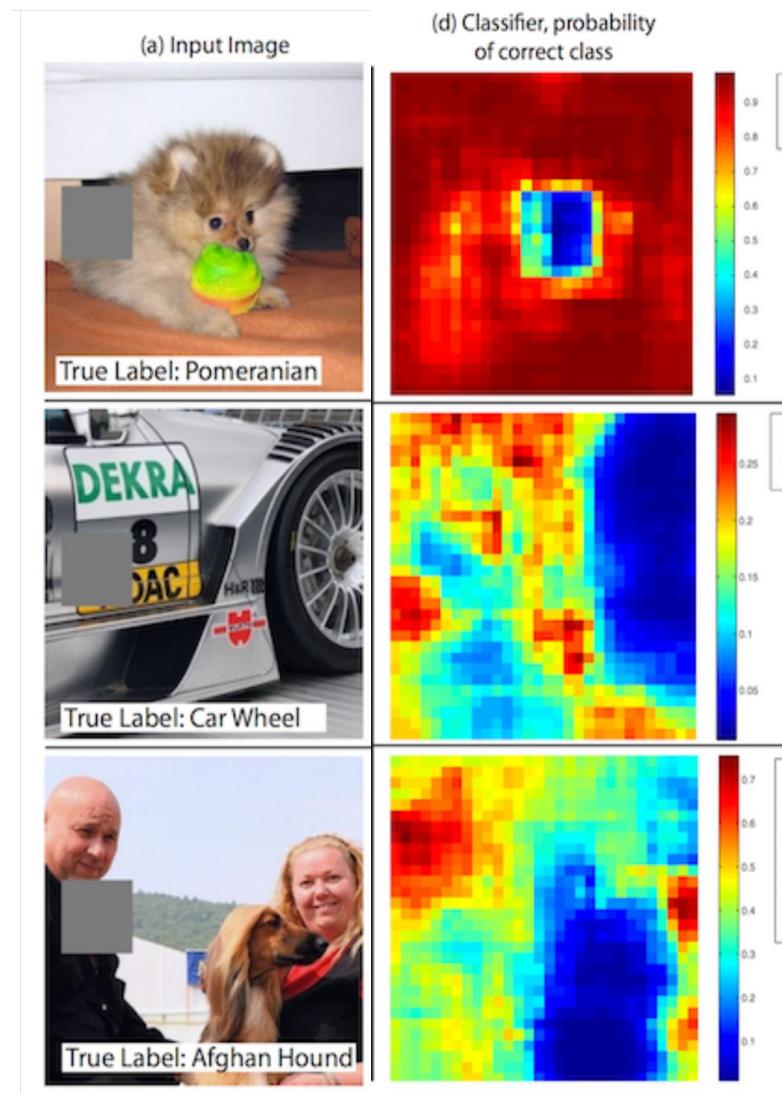
[Selvaraju, 2016, Grad-CAM: Visual Explanations from Deep Networks via Gradient-based Localization]

Guided GRAD-CAM



Occlusion-based

- Modify the input image by masking patches (with varying patch size) and record the probability of correctly classifying the class
- When the prediction drops, the pixels are maximally relevant
- Simple to implement but computationally expensive



Summary

- While training:
 - ◆ Monitor the loss and metrics on training and validation set
 - ◆ Use Tensorboard to diagnose failure to convergence
- After training:
 - ◆ Use tSNE and filter visualization to get an overall intuition of what the model has learnt
 - ◆ Inspect instance-based explanations (e.g., GradCAM)
 - ◆ Do not focus only on cases that are classified correctly: deep neural networks are prone to taking shortcuts
 - ◆ Embedding visualization works well with any type of input
 - ◆ Some types of visualization are well suited to images, possibly text
- Check available notebooks → 

Exercise

- For one of the networks trained in the previous labs (e.g., CNN trained on FashionMNIST):
 - ◆ Visualize the filters learnt using gradient ascent
 - ◆ Extract the embeddings and project them in 2 dimensions using t-SNE
 - ◆ Visualize explanations on 10 randomly selected cases using GRAD-CAM